

Hand Pose Recognition using Deep Learning

Dr. Eko Mulyanto Yuniarno

November 2024

Contents

1	Introduction	1
1.1	Importance of Hand Pose Recognition	1
1.2	Challenges in Hand Pose Recognition	2
1.3	Hand Pose Recognition Implementation in Medical	2
2	Fundamentals of Hand Pose Recognition with MediaPipe	3
2.1	Understanding Pose Recognition	3
2.2	MediaPipe	3
2.3	Real-Time Hand Tracking	3
2.4	Landmark Detection	4
2.5	Model architecture	4
2.6	Pose Estimation in 3D	5
2.7	Multi-Hand Support	5
3	Hand Pose Recognition using Deep Learning	7
3.1	Dataset Acquisition	7
3.2	Pose Estimation	8
3.2.1	Import required Python libraries	9
3.2.2	MediaPipe Initialization	9
3.2.3	Webcam Access and Directory Setup	9
3.2.4	Landmark Extraction	10
3.3	Geometric Feature Extraction	10
3.4	Model Architecture	10
3.5	Pose Classification	10
3.6	Pose Inference	10
4	Fundamentals of Hand Pose Recognition	11
4.1	Understanding Pose Recognition	11
4.2	Real-Time Hand Tracking	11
4.3	Landmark Detection	12
4.4	Model architecture	12

4.5	Pose Estimation in 3D	13
4.6	Multi-Hand Support	13

Chapter 1

Introduction

Hand pose are one of the most common communication methods in daily human life. Interaction modalities of user interfaces play a dominant role in the relationship between people and computer technology. The way users interact with interfaces has undergone a significant transformation, with most of the population now using touch devices. In today's technological advancements, human-computer interfaces (HCI) are highly appealing, mainly because they aim to enhance human lifestyles. Alongside these developments, technologies that facilitate interaction with these devices are also required. Initially, interactions were conducted using a mouse and keyboard with computers. However, with the rise of ubiquitous computing, gestures have become more prevalent—for example, smartphone interactions often involve hand gestures. The human body provides a wide range of poses that can be used as computer input. Images captured by cameras exhibit a vast amount of variation. This occurs because images are spatial data where each pixel represents a color at a specific coordinate. For pose classification, many images are needed for training to account for variations in scale, position, and hand orientation within the images. The desired outcome is a feature invariant to scale, position, and orientation, ensuring accurate and robust classification.

1.1 Importance of Hand Pose Recognition

Hand pose recognition is pivotal in advancing human-computer interaction (HCI), enabling intuitive communication between humans and machines. It bridges the gap between natural human gestures and machine understanding, opening new avenues in virtual reality, robotics, and accessibility technologies. For instance, hand gesture recognition is essential for sign language translation, empowering people with hearing impairments to communicate

seamlessly with others. Moreover, the ability to accurately recognize hand poses enhances precision in applications such as gaming, remote robotic control, and augmented reality, where fine motor movements dictate user experience. As technology continues to evolve, the importance of hand pose recognition grows, particularly in creating inclusive, adaptive, and user-friendly systems.

1.2 Challenges in Hand Pose Recognition

Despite its significance, hand pose recognition presents numerous challenges, primarily due to the variability and complexity of human hands. Hands have intricate structures with multiple degrees of freedom, leading to a vast range of poses that can be difficult to capture and interpret accurately. Variations in lighting, occlusions caused by overlapping fingers, and differences in hand shapes across individuals add further complexity. Real-time processing demands also pose a challenge, requiring high computational efficiency without sacrificing accuracy. Additionally, datasets used for training hand pose recognition models often need more diversity, limiting their generalizability in real-world scenarios. Addressing these challenges requires robust algorithms capable of handling variations and noise while maintaining computational efficiency.

1.3 Hand Pose Recognition Implementation in Medical

Hand pose recognition is revolutionizing patient care and therapy methods in the medical field. Surgeons can utilize gesture-based systems to interact with medical imaging during procedures, eliminating the need for physical contact with equipment and ensuring sterility in the operating room. In rehabilitation, hand pose recognition monitors and guides patients recovering from injuries or surgeries involving fine motor skills. For example, it enables therapists to track real-time progress and customize exercises based on a patient's movements. Moreover, hand pose recognition aids in developing assistive devices for individuals with motor impairments, enhancing their ability to perform daily tasks independently. This technology's application in healthcare improves the precision of treatments and fosters patient engagement and empowerment.

Chapter 2

Fundamentals of Hand Pose Recognition with MediaPipe

2.1 Understanding Pose Recognition

Pose recognition is analyzing and interpreting objects' physical positions and movements, particularly the human body or hands, to identify specific gestures or actions. In the context of hand pose recognition, the objective is to track and understand the intricate movements and orientations of the hand in a given space. This involves detecting the hand's position, identifying key landmarks, and classifying the overall hand configuration. Pose recognition systems combine advanced computer vision techniques with machine learning models to make predictions, offering applications in sign language interpretation, virtual reality, and human-computer interaction.

2.2 MediaPipe

MediaPipe is a cross-platform framework by Google that offers efficient pipelines for machine learning and computer vision tasks, such as pose estimation, hand tracking, and facial landmark detection. It supports mobile devices, web, and desktop environments with pre-trained models.

2.3 Real-Time Hand Tracking

Real-time hand tracking is the ability to detect and monitor hand movements instantaneously. This process relies on high-speed algorithms that process visual data from cameras or sensors to pinpoint the location and orientation of

a hand in every frame of a video stream. A significant focus of real-time hand tracking is minimizing latency to ensure the system’s responsiveness. Techniques like region-of-interest optimization, GPU acceleration, and lightweight neural networks contribute to achieving this. Real-time tracking is critical for applications such as augmented reality (AR), gaming, and touchless interfaces where seamless interaction is essential.

2.4 Landmark Detection

Landmark detection is a core element of hand pose recognition. It involves identifying key points or nodes on a hand, such as joints, knuckles, or fingertips, collectively defining its structure. These landmarks are typically represented in 2D or 3D coordinates, forming the foundation for analyzing hand gestures and poses. Advanced models like Mediapipe Hand and OpenPose leverage machine learning to perform this task efficiently. Landmark detection is critical for understanding hand dynamics, as it provides the data needed for gesture classification, tracking movements, and reconstructing hand positions in 3D space. The hand landmark can be seen in Figure 4.1.

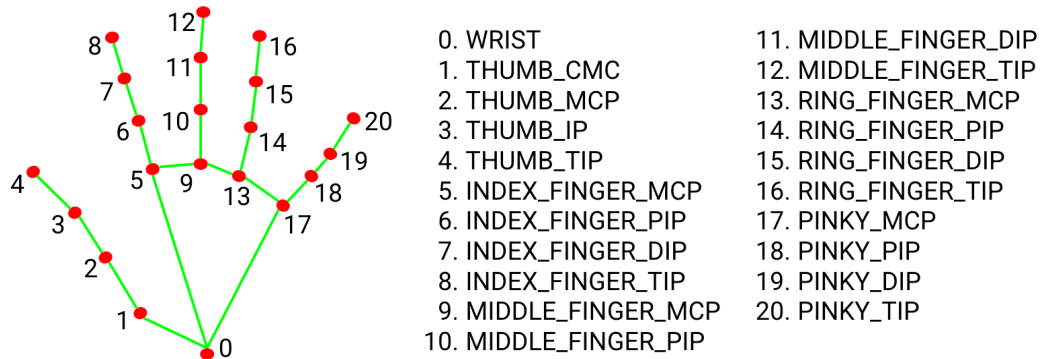


Figure 2.1: Example of Hand Landmark obtained from Mediapipe

2.5 Model architecture

The model architecture for hand pose recognition typically combines convolutional neural networks (CNNs) for feature extraction with regression or classification layers for predicting hand poses. State-of-the-art models often utilize encoder-decoder frameworks to process input images and generate

heatmaps or landmark coordinates. Mediapipe, for instance, uses a pipeline that first detects the hand region and then applies a specialized model to predict precise landmarks. Lightweight and efficient architectures are crucial to ensure real-time performance without compromising accuracy. These models are trained on large datasets to generalize across different hand shapes, orientations, and lighting conditions.

2.6 Pose Estimation in 3D

Pose estimation in 3D extends the capabilities of 2D models by incorporating depth information to reconstruct hand poses in a three-dimensional space. This process often integrates data from depth sensors or multi-view camera setups to triangulate landmark positions. 3D pose estimation provides more accurate representations of hand gestures, especially in scenarios involving rotations or occlusions. It is beneficial in applications like robotics, where precise hand positioning is required, or virtual reality, where realistic interactions with virtual objects depend on spatial accuracy. Each landmark consists of the following:

1. x and y : Landmark coordinates normalized to $[0.0, 1.0]$ by the image width and height respectively.
2. z : Represents the landmark depth with the depth at the midpoint of hips being the origin, and the smaller the value the closer the landmark is to the camera. The magnitude of z uses roughly the same scale as x .
3. visibility: A value in $[0.0, 1.0]$ indicating the likelihood of the landmark being visible (present and not occluded) in the image.

2.7 Multi-Hand Support

Multi-hand support is the capability of a hand pose recognition system to track and analyze multiple hands within the same frame simultaneously. This requires the model to distinguish between hands, assign unique identifiers, and accurately detect the landmarks for each hand. Challenges in multi-hand support include managing occlusions (where one hand partially covers the other) and handling varying hand orientations. Robust systems, like Mediapipe, employ efficient region segmentation and tracking algorithms to maintain consistent performance in multi-hand scenarios. Multi-hand support is essential for collaborative applications like multi-user gaming or shared virtual workspaces.

Chapter 3

Hand Pose Recognition using Deep Learning

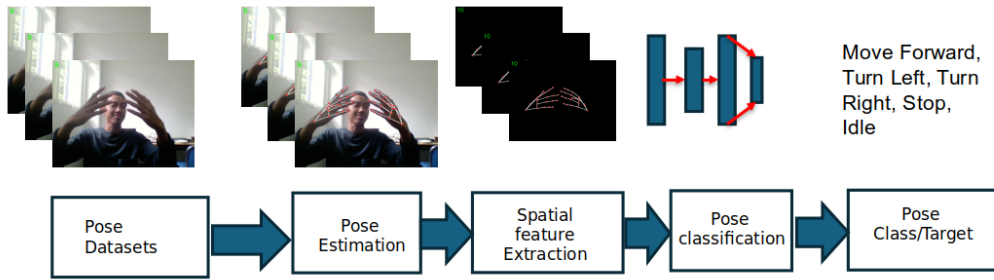


Figure 3.1: Hand Recognition Pipeline

3.1 Dataset Acquisition

This code provides a systematic way to acquire a dataset for hand pose recognition using MediaPipe and OpenCV. It captures hand landmarks and saves the corresponding frames from a webcam feed, enabling the creation of a labeled dataset for tasks like training gesture recognition models.

The process begins with the initialization and setup of necessary components. MediaPipe's Hands model is used for detecting and tracking hand landmarks, while OpenCV handles video capture and frame display. Command-line arguments allow customization of parameters such as the dataset path, label name, frame saving rate, maximum frames to save, and a delay before data collection starts. The script ensures that a directory for the dataset is created, organizing the captured data based on the provided label name.

Hand detection and landmark extraction are key steps in the process. The Mediapipe Hands model detects 21 landmarks on each hand, including knuckles and fingertips, and supports tracking up to two hands simultaneously. For every frame captured, the script converts it to RGB format and processes it using Mediapipe. Landmark coordinates are normalized relative to specific reference points, such as the wrist and a finger joint, making the data invariant to variations in hand size and position. The normalized landmarks for both hands are then concatenated into a single array, representing the frame's hand pose data.

3.2 Pose Estimation

The `'create_dataset'` function is a crucial part of the pose estimation process. It facilitates the acquisition of hand pose data by capturing frames from a webcam feed, extracting hand landmarks using Mediapipe's 'Hands' solution, and organizing the data into labeled directories. This function serves as the foundation for building datasets tailored to pose estimation tasks, enabling further applications like gesture recognition and motion analysis.

Pose estimation involves detecting and interpreting key hand landmarks in real-time. The function initializes Mediapipe's 'Hands' module, which is capable of identifying 21 specific landmarks on each hand, such as fingertips, joints, and the wrist. These landmarks represent the structural configuration of the hand and are essential for understanding its pose. By converting the captured video frames to RGB format and processing them through Mediapipe, the function extracts these landmarks with high precision.

To normalize the landmark data, the function computes positions relative to key reference points—namely, the wrist and a finger joint. This normalization step ensures that the data remains invariant to hand size and camera perspective, enhancing its generalizability for machine learning models. The landmark coordinates for both hands are flattened into a single array, representing the spatial configuration of the hands in each frame.

Beyond pose estimation, the `'create_dataset'` function also incorporates real-time feedback and organization. Frames are captured and saved at user-defined intervals, allowing control over the frame rate and total number of frames collected. Each frame is stored in a labeled directory, organized by the name of the gesture or pose being recorded. The overlay of the detected landmarks on the video feed, along with a live counter of saved frames, enhances the usability of the function by providing visual confirmation and progress tracking.

This function exemplifies the integration of pose estimation techniques

with practical data collection workflows. It not only extracts meaningful pose information but also structures the data in a format ready for subsequent analysis and model training. This makes it a powerful tool for researchers and developers working on applications in gesture recognition, human-computer interaction, and beyond.

3.2.1 Import required Python libraries

```
import cv2
import mediapipe as mp
import numpy as np
from datetime import datetime
import os
import time
import argparse
```

3.2.2 MediaPipe Initialization

```
mp_hands = mp.solutions.hands
hands = mp_hands.Hands(static_image_mode=False ,
max_num_hands=2,
min_detection_confidence=0.5,
min_tracking_confidence=0.5)
```

This part initializes the MediaPipe Hands solution, which is responsible for detecting hand landmarks in video frames. The configuration allows real-time processing of up to two hands per frame, with parameters controlling the confidence thresholds for detection and tracking. This setup ensures accurate and stable landmark recognition during data collection.

3.2.3 Webcam Access and Directory Setup

```
cap = cv2.VideoCapture(0)
save_dir = os.path.join(direktori_path , nama_label)
if not os.path.exists(save_dir):
    os.makedirs(save_dir)
```

The webcam is opened using OpenCV's VideoCapture, which serves as the data source for real-time frame capture. The save_dir is created dynamically based on the dataset path and label name, ensuring an organized directory structure for storing the collected data.

3.2.4 Landmark Extraction

```
def ekstraksi_fitur(frame):
    frame_rgb = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
    results = hands.process(frame_rgb)
    height, width, _ = frame.shape
    ...
    # Normalization and concatenation of landmarks
    return np.concatenate((left_hand_landmarks.flatten(), right_hand_landmarks.flatten()))
```

This function extracts and processes hand landmarks from each frame:

1. Conversion to RGB: Mediapipe requires frames in RGB format for processing.
2. Landmark Processing: If hand landmarks are detected, their coordinates are normalized relative to reference points (e.g., the wrist and a finger joint) to ensure invariance to hand size and position.
3. Concatenation: The landmarks for both hands are flattened and concatenated into a single array, representing the full pose information for the frame.

3.3 Geometric Feature Extraction

3.4 Model Architecture

3.5 Pose Classification

3.6 Pose Inference

Chapter 4

Fundamentals of Hand Pose Recognition

4.1 Understanding Pose Recognition

Pose recognition is analyzing and interpreting objects' physical positions and movements, particularly the human body or hands, to identify specific gestures or actions. In the context of hand pose recognition, the objective is to track and understand the intricate movements and orientations of the hand in a given space. This involves detecting the hand's position, identifying key landmarks, and classifying the overall hand configuration. Pose recognition systems combine advanced computer vision techniques with machine learning models to make predictions, offering applications in sign language interpretation, virtual reality, and human-computer interaction.

4.2 Real-Time Hand Tracking

Real-time hand tracking is the ability to detect and monitor hand movements instantaneously. This process relies on high-speed algorithms that process visual data from cameras or sensors to pinpoint the location and orientation of a hand in every frame of a video stream. A significant focus of real-time hand tracking is minimizing latency to ensure the system's responsiveness. Techniques like region-of-interest optimization, GPU acceleration, and lightweight neural networks contribute to achieving this. Real-time tracking is critical for applications such as augmented reality (AR), gaming, and touchless interfaces where seamless interaction is essential.

4.3 Landmark Detection

Landmark detection is a core element of hand pose recognition. It involves identifying key points or nodes on a hand, such as joints, knuckles, or fingertips, collectively defining its structure. These landmarks are typically represented in 2D or 3D coordinates, forming the foundation for analyzing hand gestures and poses. Advanced models like Mediapipe Hand and OpenPose leverage machine learning to perform this task efficiently. Landmark detection is critical for understanding hand dynamics, as it provides the data needed for gesture classification, tracking movements, and reconstructing hand positions in 3D space. The hand landmark can be seen in Figure 4.1.




Figure 4.1: Example of Hand Landmark obtained from Mediapipe

4.4 Model architecture

The model architecture for hand pose recognition typically combines convolutional neural networks (CNNs) for feature extraction with regression or classification layers for predicting hand poses. State-of-the-art models often utilize encoder-decoder frameworks to process input images and generate heatmaps or landmark coordinates. Mediapipe, for instance, uses a pipeline that first detects the hand region and then applies a specialized model to predict precise landmarks. Lightweight and efficient architectures are crucial to ensure real-time performance without compromising accuracy. These models are trained on large datasets to generalize across different hand shapes, orientations, and lighting conditions.

4.5 Pose Estimation in 3D

Pose estimation in 3D extends the capabilities of 2D models by incorporating depth information to reconstruct hand poses in a three-dimensional space. This process often integrates data from depth sensors or multi-view camera setups to triangulate landmark positions. 3D pose estimation provides more accurate representations of hand gestures, especially in scenarios involving rotations or occlusions. It is beneficial in applications like robotics, where precise hand positioning is required, or virtual reality, where realistic interactions with virtual objects depend on spatial accuracy. Each landmark consists of the following:

1. x and y : Landmark coordinates normalized to $[0.0, 1.0]$ by the image width and height respectively.
2. z : Represents the landmark depth with the depth at the midpoint of hips being the origin, and the smaller the value the closer the landmark is to the camera. The magnitude of z uses roughly the same scale as x .
3. visibility: A value in $[0.0, 1.0]$ indicating the likelihood of the landmark being visible (present and not occluded) in the image.

4.6 Multi-Hand Support

Multi-hand support is the capability of a hand pose recognition system to track and analyze multiple hands within the same frame simultaneously. This requires the model to distinguish between hands, assign unique identifiers, and accurately detect the landmarks for each hand. Challenges in multi-hand support include managing occlusions (where one hand partially covers the other) and handling varying hand orientations. Robust systems, like Mediapipe, employ efficient region segmentation and tracking algorithms to maintain consistent performance in multi-hand scenarios. Multi-hand support is essential for collaborative applications like multi-user gaming or shared virtual workspaces.

